

Implementing Bayes by Backprop

Juho Lee

Grad school of AI

Samsung AI expert course

dense layer in tensorflow

```
def dense(x, num_units, name='dense', activation=None):  
    with tf.variable_scope(name, reuse=tf.AUTO_REUSE):  
        W = tf.get_variable('W', [x.shape[1], num_units])  
        b = tf.get_variable('b', [num_units],  
                             initializer=tf.zeros_initializer())  
        x = tf.matmul(x, W) + b  
        if activation == 'relu':  
            x = tf.nn.relu(x)  
        return x
```

$$\mathbf{y} = \mathbf{x}\mathbf{W} + \mathbf{b}$$

$\mathbb{R}^{N \times d_{\text{in}}}$ $\mathbb{R}^{d_{\text{in}} \times d_{\text{out}}}$ $\mathbb{R}^{d_{\text{out}}}$

Bayes Dense?

- Keep the distribution of W and b
- The distributions? - Gaussians
- To define Gaussians, we need means and covariances.

bayes_dense layers in tensorflow

```
def bayes_dense(x, num_units, name='dense', gamma=1.0, activation=None):  
    with tf.variable_scope(name, reuse=tf.AUTO_REUSE):  
        W_mu = tf.get_variable('W_mu', [x.shape[1], num_units])  
        W_rho = tf.nn.softplus(  
            tf.get_variable('W_rho', [x.shape[1], num_units],  
                            initializer=tf.random_uniform_initializer(-3., -2.)))  
        b_mu = tf.get_variable('b_mu', [num_units],  
                                initializer=tf.zeros_initializer())  
        b_rho = tf.nn.softplus(  
            tf.get_variable('b_rho', [num_units],  
                            initializer=tf.random_uniform_initializer(-3., -2.)))
```

$$q(\mathbf{W}) = \mathcal{N}(\boldsymbol{\mu}_{\mathbf{W}}, \boldsymbol{\rho}_{\mathbf{W}}^2) = \prod_{i,j} \mathcal{N}(\mu_{w_{ij}}, \rho_{w_{ij}}^2)$$

$$q(\mathbf{b}) = \mathcal{N}(\boldsymbol{\mu}_{\mathbf{b}}, \boldsymbol{\rho}_{\mathbf{b}}^2) = \prod_j \mathcal{N}(\mu_{b_j}, \rho_{b_j}^2)$$

Computing outputs in bayes_dense

```
# sample
W = W_mu + W_rho * tf.random.normal(W_mu.shape)
b = b_mu + b_rho * tf.random.normal(b_mu.shape)

x = tf.matmul(x, W) + b
if activation == 'relu':
    x = tf.nn.relu(x)
```

$$\mathbf{W} = \boldsymbol{\mu}_{\mathbf{W}} + \boldsymbol{\varepsilon} \odot \boldsymbol{\rho}_{\mathbf{W}}, \quad \boldsymbol{\varepsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$$

$$\mathbf{b} = \boldsymbol{\mu}_{\mathbf{b}} + \boldsymbol{\varepsilon} \odot \boldsymbol{\rho}_{\mathbf{b}}, \quad \boldsymbol{\varepsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$$

KL-divergence

```
# kl divergence
kld_W = tf.reduce_sum(kl_divergence(Normal(W_mu, W_rho), Normal(0., gamma)))
kld_b = tf.reduce_sum(kl_divergence(Normal(b_mu, b_rho), Normal(0., gamma)))
kld = kld_W + kld_b
```

$$p(\mathbf{W}) = \mathcal{N}(\mathbf{0}, \gamma \mathbf{I}) \quad p(\mathbf{b}) = \mathcal{N}(\mathbf{0}, \gamma \mathbf{I})$$

$$\begin{aligned} \text{KL}[q(\mathbf{W})q(\mathbf{b})||p(\mathbf{W})p(\mathbf{b})] = \\ \text{KL}[q(\mathbf{W})||p(\mathbf{W})] + \text{KL}[q(\mathbf{b})||p(\mathbf{b})] \end{aligned}$$